

Double Oracle Algorithm for Computing Equilibria in Continuous Games

Lukáš Adam¹, Rostislav Horčík¹, Tomáš Kasl¹, Tomáš Kroupa¹

¹Department of Computer Science, Faculty of Electrical Engineering, Czech Technical University in Prague
{ adamluk3, xhorcik, kasltoma, tomas.kroupa }@fel.cvut.cz

Abstract

Many efficient algorithms have been designed to recover Nash equilibria of various classes of finite games. Special classes of continuous games with infinite strategy spaces, such as polynomial games, can be solved by semidefinite programming. In general, however, continuous games are not directly amenable to computational procedures. In this contribution, we develop an iterative strategy generation technique for finding a Nash equilibrium in a whole class of continuous two-person zero-sum games with compact strategy sets. The procedure, which is called the double oracle algorithm, has been successfully applied to large finite games in the past. We prove the convergence of the double oracle algorithm to a Nash equilibrium. Moreover, the algorithm is guaranteed to recover an approximate equilibrium in finitely-many steps. Our numerical experiments show that it outperforms fictitious play on several examples of games appearing in the literature. In particular, we provide a detailed analysis of experiments with a version of the continuous Colonel Blotto game.

1 Introduction

Action spaces of games appearing in AI applications are often prohibitively large. Consequently, one has to strive for efficiently computable approximations of equilibria, possibly with provable bounds on convergence rates (Gilpin, Peña, and Sandholm 2012). A number of algorithms applied in AI like fictitious play (Brown 1951), the double oracle algorithm (McMahan, Gordon, and Blum 2003) or the policy-space response oracle (Lanctot et al. 2017; Muller et al. 2019) overcome the problem with the cardinality by selecting ‘good’ strategies iteratively. The selection process is usually based on an approximation of the best response. In a nutshell, the recent advances in algorithmic game theory has led to the development of algorithms for (approximately) solving extremely large finite games, such as variants of poker (Moravčík et al. 2017; Brown and Sandholm 2019) or multidimensional resource allocation problems (Behnezhad et al. 2017).

Completely new problems arise from considering games with infinite strategy spaces, in which the strategies are vectors of real numbers corresponding to physical parameters (Archibald and Shoham 2009) or to the setting of classifiers

(Yasodharan and Loiseau 2019). The first theoretical obstacle is that the existence of mixed strategy equilibria is guaranteed only for infinite games whose utility functions satisfy additional conditions (Glicksberg 1952; Fan 1952). On top of that, some well understood classes of infinite games possess only optimal strategies whose supports are uncountable; see (Roberson 2006) for an in-depth discussion of infinite Colonel Blotto games.

Computational procedures for finding (approximate) equilibria of infinite games exist for rather special kinds of utility functions. Two-person zero-sum polynomial games are solvable by semidefinite programming; see (Parrilo 2006; Laraki and Lasserre 2012). Approximate equilibria of separable games can be computed under additional assumptions (Stein, Ozdaglar, and Parrilo 2008). However, games appearing in applications are rarely of the form above and a detailed analysis of their properties is inevitable; see (Yasodharan and Loiseau 2019) for an application in adversarial machine learning. Some authors develop approximations of best response by neural nets (Kamra et al. 2018, 2019). One of the important iterative procedures for finite games, Brown-Robinson learning process known as fictitious play (Brown 1951; Robinson 1951), has been recently applied to infinite games (Ganzfried 2020). However, the dynamics of best response strategies generated by fictitious play was analyzed only in special cases; cf. (Hofbauer and Sorin 2006; Perkins and Leslie 2014). To the best of our knowledge, not much is known about the convergence of fictitious play for general zero-sum continuous games as defined below.

This paper deals with continuous games, which we define as two-person zero-sum games with continuous utility functions over compact strategy sets. We extend the double oracle algorithm (McMahan, Gordon, and Blum 2003) to such games. This algorithm is an iterative strategy generation technique based on (i) the solution of subgames by LP solvers and (ii) the expansion of subgames’ strategy sets using the best response strategies obtained thus far. Our main result is the convergence of this algorithm for any continuous game (Theorem 3.1). The numerical experiments in Section 4 show that the double oracle algorithm converges faster than fictitious play on several examples (polynomial game, Townsend function, and a version of the Colonel Blotto game). The repository with our experiments’ codes is <https://github.com/sadda/DoubleOracle>.

2 Basic Notions

This section summarizes basic notions and results related to continuous zero-sum games and their equilibria; see (Karlin 1959) or (Stein, Ozdaglar, and Parrilo 2008) for details.

Continuous Games

Player 1 and Player 2 select strategies from nonempty compact sets $X \subseteq \mathbb{R}^m$ and $Y \subseteq \mathbb{R}^n$, respectively. The utility function of Player 1 is a continuous function $u: X \times Y \rightarrow \mathbb{R}$. The utility function of Player 2 is $-u$. The triple $\mathcal{G} = (X, Y, u)$ is called a *continuous game*. Note that some authors use the term ‘continuous game’ in a somewhat different sense allowing utility functions to be discontinuous functions over metric spaces of strategies.

A continuous game $\mathcal{G} = (X, Y, u)$ is (i) *finite* if both X and Y are finite, and (ii) *infinite* if X or Y is infinite. We will need the notion of subgame. When $X' \subseteq X$ and $Y' \subseteq Y$ are nonempty compact sets, we define the *subgame* $\mathcal{G}' = (X', Y', u)$ of \mathcal{G} by the restriction of u to $X' \times Y'$, which is denoted by the same letter.

The concept of mixed strategy in continuous games should allow every player to randomize with respect to any probability measure on the corresponding strategy set. We will spell out the definitions related to mixed strategies only for Player 1. Their counterparts for Player 2 are completely analogous. A *mixed strategy* of Player 1 is a Borel probability measure p over X . The set of all mixed strategies of Player 1 is denoted by Δ_X . The *support* of a mixed strategy $p \in \Delta_X$ is the set

$$\text{spt } p := \bigcap \{K \subseteq X \mid K \text{ compact, } p(K) = 1\}.$$

Every mixed strategy $p \in \Delta_X$ can be classified as one of the following types depending on the size of its support.

1. *Pure strategy* p . This means that $\text{spt } p = \{x\}$ for some $x \in X$. Equivalently, p is equal to Dirac measure δ_x .
2. *Finitely-supported mixed strategy* p . The support $\text{spt } p$ is finite. Hence, p can be written as a convex combination

$$p = \sum_{x \in \text{spt } p} p(x) \cdot \delta_x.$$

3. *Mixed strategy* p with *infinite support* $\text{spt } p$.

Put $\Delta := \Delta_X \times \Delta_Y$. If players implement a mixed strategy profile $(p, q) \in \Delta$, the expected utility of Player 1 is

$$U(p, q) := \int_{X \times Y} u(x, y) \, d(p \times q). \quad (1)$$

This yields a function $U: \Delta \rightarrow \mathbb{R}$, which can be effectively evaluated in important special cases. For example, when both $\text{spt } p$ and $\text{spt } q$ are finite,

$$U(p, q) = \sum_{x \in \text{spt } p} \sum_{y \in \text{spt } q} p(x) \cdot q(y) \cdot u(x, y).$$

If Player 1 employs a pure strategy given by $x \in X$ and Player 2 uses a mixed strategy $q \in \Delta_Y$, we will use the short notation $U(x, q) := U(\delta_x, q)$.

Equilibria in Continuous Games

A mixed strategy profile $(p^*, q^*) \in \Delta$ is an *equilibrium* in a continuous game \mathcal{G} if

$$U(p, q^*) \leq U(p^*, q^*) \leq U(p^*, q) \quad (2)$$

holds for all $(p, q) \in \Delta$. By Glicksberg’s theorem (Glicksberg 1952), every continuous game has an equilibrium. Define the *lower/upper value* of \mathcal{G} by

$$\begin{aligned} \underline{v}(\mathcal{G}) &:= \max_{p \in \Delta_X} \min_{q \in \Delta_Y} U(p, q) \quad \text{and} \\ \bar{v}(\mathcal{G}) &:= \min_{q \in \Delta_Y} \max_{p \in \Delta_X} U(p, q). \end{aligned}$$

Proposition 2.1 gives several conditions for equilibrium, which will be used throughout the paper without further references. Its proof is omitted since it is completely analogous to the case of finite games.

Proposition 2.1. *Let $\mathcal{G} = (X, Y, u)$ be a continuous game and $(p^*, q^*) \in \Delta$. The following assertions are equivalent.*

1. *The strategy profile (p^*, q^*) is an equilibrium.*
2. *$U(x, q^*) \leq U(p^*, q^*) \leq U(p^*, y)$ for all $(x, y) \in X \times Y$.*
3. *$\min_{y \in Y} U(p^*, y) = \underline{v}(\mathcal{G})$ and $\max_{x \in X} U(x, q^*) = \bar{v}(\mathcal{G})$.*
4. *$\underline{v}(\mathcal{G}) = U(p^*, q^*) = \bar{v}(\mathcal{G})$.*

Hence, the equality $\underline{v}(\mathcal{G}) = \bar{v}(\mathcal{G})$ holds for every continuous game \mathcal{G} , and $v(\mathcal{G}) := \underline{v}(\mathcal{G})$ is called the *value* of \mathcal{G} .

Bounds on the size of supports of equilibrium strategies are known only for particular classes of continuous games, such as the class of separable games (Stein, Ozdaglar, and Parrilo 2008). There are examples of games whose equilibria are almost any sets of finitely-supported mixed strategies (Rehbeck 2018). Moreover, some continuous games possess only equilibria with uncountable supports (Roberson 2006).

In many applications it is enough to find an ϵ -equilibrium (p^*, q^*) for some $\epsilon \geq 0$, that is,

$$U(p, q^*) - \epsilon \leq U(p^*, q^*) \leq U(p^*, q) + \epsilon \quad (3)$$

for all $(p, q) \in \Delta$. Note that this is a natural extension of (2). According to Proposition 2.2, whose proof is in Appendix B, we can always recover an approximate equilibrium (p^*, q^*) with finite supports and such that $U(p^*, q^*)$ is arbitrarily close to the value of game $v(\mathcal{G})$.

Proposition 2.2. *Let \mathcal{G} be an arbitrary continuous game. Then for every $\epsilon > 0$:*

- *There exists an ϵ -equilibrium (p^*, q^*) of \mathcal{G} such that both $\text{spt } p^*$ and $\text{spt } q^*$ are finite.*
- *Every ϵ -equilibrium (p^*, q^*) of \mathcal{G} satisfies the inequality $|U(p^*, q^*) - v(\mathcal{G})| \leq \epsilon$.*

3 Double Oracle Algorithm

The double oracle algorithm uses the notion of best response strategies. For every mixed strategy $q \in \Delta_Y$ of Player 2, the *best response set* of Player 1 is

$$\beta_1(q) := \left\{ x \in X \mid U(x, q) = \max_{x' \in X} U(x', q) \right\}.$$

Analogously, for any $p \in \Delta_X$, put

$$\beta_2(p) := \left\{ y \in Y \mid U(p, y) = \min_{y' \in Y} U(p, y') \right\}.$$

Note that best response strategies are defined to be pure, without any loss of generality; see Proposition A.2. Moreover, by compactness and continuity, $\beta_1(q)$ and $\beta_2(p)$ are always nonempty compact sets.

The idea of the double oracle algorithm (Algorithm 3.1) applied to a continuous game $\mathcal{G} = (X, Y, u)$ is simple. In every iteration, finite strategy sets X_i and Y_i are determined and some equilibrium (p_i^*, q_i^*) of the finite subgame (X_i, Y_i, u) is found by the standard linear programming methods. The best responses x_{i+1} and y_{i+1} to q_i^* and p_i^* , respectively, are recovered, and added to the strategy sets. This is repeated until a terminating condition is satisfied. The resulting strategy profile is guaranteed to be an ϵ -equilibrium.

Algorithm 3.1 Double Oracle Algorithm

Input: Continuous game $\mathcal{G} = (X, Y, u)$, nonempty finite subsets $X_1 \subseteq X$, $Y_1 \subseteq Y$, and $\epsilon \geq 0$

- 1: Let $i := 0$
- 2: **repeat**
- 3: Increase i by one
- 4: Find an equilibrium (p_i^*, q_i^*) of subgame (X_i, Y_i, u)
- 5: Find some $x_{i+1} \in \beta_1(q_i^*)$ and $y_{i+1} \in \beta_2(p_i^*)$
- 6: Let $X_{i+1} := X_i \cup \{x_{i+1}\}$ and $Y_{i+1} := Y_i \cup \{y_{i+1}\}$
- 7: Let $\underline{v}_i := U(p_i^*, y_{i+1})$ and $\bar{v}_i := U(x_{i+1}, q_i^*)$
- 8: **until** $\bar{v}_i - \underline{v}_i \leq \epsilon$

Output: ϵ -equilibrium (p_i^*, q_i^*) of game \mathcal{G}

We now perform a simple analysis of the algorithm. Since

$$U(p_i^*, q_i^*) = \max_{x \in X_i} U(x, q_i^*) \leq \max_{x \in X} U(x, q_i^*) = \bar{v}_i$$

and similarly for the lower bound, we have

$$\underline{v}_i \leq U(p_i^*, q_i^*) \leq \bar{v}_i. \quad (4)$$

Lemma B.2 states that the same bounds hold even for the value of game \mathcal{G} :

$$\underline{v}_i \leq v(\mathcal{G}) \leq \bar{v}_i.$$

The usual stopping condition of the double oracle algorithm for finite games is $X_{i+1} = X_i$ and $Y_{i+1} = Y_i$. Herein we chose the terminating condition $\bar{v}_i - \underline{v}_i \leq \epsilon$ for two reasons:

- It is more general. Indeed, Lemma B.1 states that if $X_{i+1} = X_i$ and $Y_{i+1} = Y_i$, then $\bar{v}_i - \underline{v}_i = 0$.
- It provides an estimate for the quality of approximate equilibrium. Formula (4) implies $\bar{v}_i - \underline{v}_i \geq 0$ and Theorem 3.1 states that (p_i^*, q_i^*) is an $(\bar{v}_i - \underline{v}_i)$ -equilibrium. Then Proposition 2.2 guarantees that $v(\mathcal{G})$ is known precisely up to $(\bar{v}_i - \underline{v}_i)$.

The main result of this manuscript is the convergence of the double oracle algorithm. In fact our result generalizes the result about convergence of the double oracle algorithm for finite games; see (McMahan, Gordon, and Blum 2003). For finite games, we neglect the case of $\epsilon > 0$ since the algorithm is known to converge to an equilibrium for $\epsilon = 0$ in finitely many steps.

Theorem 3.1. Let $\mathcal{G} = (X, Y, u)$ be a continuous game.

1. If \mathcal{G} is a finite game and $\epsilon = 0$, Algorithm 3.1 converges to an equilibrium in a finite number of iterations.
2. If \mathcal{G} is an infinite game and $\epsilon = 0$, every weakly convergent subsequence of Algorithm 3.1 converges to an equilibrium in a possibly infinite number of iterations. Moreover, such a weakly convergent subsequence always exist.
3. If \mathcal{G} is an infinite game and $\epsilon > 0$, Algorithm 3.1 converges to a finitely supported ϵ -equilibrium in a finite number of iterations.

Proof. We first realize that the terminating condition

$$U(x_{i+1}, q_i^*) - U(p_i^*, y_{i+1}) \leq \epsilon$$

implies

$$\begin{aligned} U(p_i^*, q_i^*) &\leq U(x_{i+1}, q_i^*) \leq U(p_i^*, y_{i+1}) + \epsilon \\ &= \min_{y' \in Y} U(p_i^*, y') + \epsilon = \min_{q \in \Delta_Y} U(p_i^*, q) + \epsilon. \end{aligned}$$

The first and the third relation above follow from the definition of the best response, the second from the terminating condition and the last from Proposition A.2. Similarly, we can show that

$$\begin{aligned} U(p_i^*, q_i^*) &\geq U(p_i^*, y_{i+1}) \geq U(x_{i+1}, q_i^*) - \epsilon \\ &= \max_{x' \in X} U(x', q_i^*) - \epsilon = \max_{p \in \Delta_X} U(p, q_i^*) - \epsilon. \end{aligned}$$

Combining these two inequalities implies that (p_i^*, q_i^*) is an ϵ -equilibrium. Note that for $\epsilon = 0$, this means that (p_i^*, q_i^*) is an equilibrium.

Item 1. If \mathcal{G} is finite, then after a finite number of iterations it must happen that $X_{i+1} = X_i$ and $Y_{i+1} = Y_i$. Lemma B.1 implies that the terminating condition of Algorithm 3.1 is satisfied with $\epsilon = 0$ and the first paragraph of this proof implies that (p_i^*, q_i^*) is an equilibrium of \mathcal{G} .

Item 2. Consider now the case of an infinite game and $\epsilon = 0$. If the double oracle algorithm terminates in a finite number of iterations, then the first paragraph implies that (p_i^*, q_i^*) is an equilibrium. In the opposite case, the algorithm produces an infinite number of iterations. Due to Proposition A.1, there is a weakly convergent subsequence which, for simplicity, will be denoted by the same indices. Therefore, $p_i^* \Rightarrow p^*$ for some p^* and $q_i^* \Rightarrow q^*$ for some q^* , where the symbol \Rightarrow denotes the weak convergence (Appendix A).

Consider any y such that $y \in Y_{i_0}$ for some i_0 . Take an arbitrary $i \geq i_0$, which implies $y \in Y_i$. Since (p_i^*, q_i^*) is an equilibrium of the subgame (X_i, Y_i, u) , we get

$$U(p_i^*, q_i^*) \leq U(p_i^*, y) \rightarrow U(p^*, y),$$

where the convergence follows from (11). Since $U(p_i^*, q_i^*) \rightarrow U(p^*, q^*)$ due to (10), this implies

$$U(p^*, q^*) \leq U(p^*, y) \quad (5)$$

for all $y \in \cup Y_i$. Since U is continuous, the previous inequality holds for all $y \in \text{cl}(\cup Y_i)$.

Fix now an arbitrary $y \in Y$. Because y_{i+1} is the best response, we get

$$U(p_i^*, y_{i+1}) \leq U(p_i^*, y) \rightarrow U(p^*, y), \quad (6)$$

where the limit holds due to (11). Since $y_{i+1} \in Y_{i+1}$ and by compactness of Y , we can select a convergent subsequence $y_i \rightarrow \hat{y}$, again without any relabelling, where $\hat{y} \in \text{cl}(\cup Y_i)$. This allows us to use (5) to obtain

$$U(p_i^*, y_{i+1}) \rightarrow U(p^*, \hat{y}) \geq U(p^*, q^*). \quad (7)$$

Combining (6) and (7) yields

$$U(p^*, q^*) \leq U(p^*, y)$$

for all $y \in Y$. Repeating the analogous arguments in the other variable yields

$$U(x, q^*) \leq U(p^*, q^*) \leq U(p^*, y)$$

for all $x \in X$ and $y \in Y$. Then Proposition 2.1 says that (p^*, q^*) is an equilibrium of \mathcal{G} .

Item 3. Consider now the case of an infinite game with $\epsilon > 0$ and realize that (6) and (7) also imply

$$\begin{aligned} U(p^*, q^*) &\leq U(p^*, \hat{y}) \leftarrow U(p_i^*, y_{i+1}) \\ &\leq U(p_i^*, q_i^*) \rightarrow U(p^*, q^*), \end{aligned}$$

which means $U(p_i^*, y_{i+1}) \rightarrow U(p^*, q^*)$. Similarly, $U(x_{i+1}, q_i^*) \rightarrow U(p^*, q^*)$ and therefore

$$\bar{v}_i - \underline{v}_i = U(x_{i+1}, q_i^*) - U(p_i^*, y_{i+1}) \rightarrow 0.$$

This states that the terminating condition will be satisfied after a finite number of iterations and the first paragraph of this proof states that (p_i^*, q_i^*) is an ϵ -equilibrium. Since only a finite number of iterations was performed and since X_1 and X_2 are finite, this implies that the supports of p_i^* and q_i^* are finite as well. \square

Since best response strategies are not unique, in general, the sequence generated by Algorithm 3.1 may fail to converge for some continuous games. Hence, it is necessary to consider a convergent subsequence of iterates in Theorem 3.1. Such a continuous game is shown in Example B.1. Another feature of the double oracle algorithm is that the sequence $\bar{v}_i - \underline{v}_i$ has nonnegative terms and converges to zero, but it is not necessarily monotone. This behavior can be demonstrated even for some finite games.

4 Numerical Experiments

We present two classes of games. The first class contains one-dimensional strategy spaces and the second class consists of certain Colonel Blotto games. The equilibrium of each finite subgame is found by solving a linear program. The best responses were computed by selecting the best point of a uniform discretization for the one-dimensional problems and by using a mixed-integer linear programming reformulation for the Colonel Blotto games. The examples were implemented in Python with solvers `scipy.optimize` and `mip`. All computations were performed on a laptop with Intel Core i5 CPU and 8GB RAM and no GPU was involved. Randomness is present only in the initialization of one-dimensional examples when a random pair of pure strategies is found.

We compare the double oracle algorithm with fictitious play. Its extension from finite to infinite games was recently formulated in (Ganzfried 2020).

One-dimensional Examples

We consider a polynomial game \mathcal{G}_1 from (Parrilo 2006) with the strategy spaces $X = Y = [-1, 1]$ and the utility function

$$u_1(x, y) = 5xy - 2x^2 - 2xy^2 - y.$$

In the equilibrium, Player 1 has the pure strategy $x^* = 0.2$ and Player 2 has the mixed strategy $q^* = 0.78\delta_1 + 0.22\delta_{-1}$. The value of game is -0.48 . Figure 1 shows the convergence of upper/lower estimates of the value of game. Note that the fictitious play is much slower to converge than the double oracle algorithm.




Figure 1: Convergence to the value of game \mathcal{G}_1

The utility function u_2 in our second example (game \mathcal{G}_2) is based on (Townsend 2014). Specifically,

$$u_2(x, y) = -\cos^2((x - 0.1)y) - x \sin(3x + y)$$

is defined on $X = [-2.25, 2.5]$ and $Y = [-2.5, 1.75]$; see Figure 2. The convergence to the value is depicted on Figure 3. Once again the double oracle algorithm converges fast, while fictitious play is rather slow to converge. In Figure 4 we show the optimal strategies of Player 1. The double oracle algorithm converged to a mixed strategy supported by four points, the fictitious play seems to reach in limit a continuous distribution whose peaks are those points. Note that the vertical axis is rescaled to account for the difference between discrete and continuous distributions.




Figure 2: Townsend function u_2




Figure 3: Convergence to the value of game \mathcal{G}_2




Figure 4: Mixed strategies in game \mathcal{G}_2

Colonel Blotto Game

We consider a continuous variant of the Colonel Blotto game. Two players simultaneously allocate forces across n battlefields. Both strategy spaces X and Y equal to

$$\left\{ \mathbf{x} := (x^1, \dots, x^n) \in \mathbb{R}^n \mid x^j \geq 0, \sum_{j=1}^n x^j = 1 \right\}.$$

The utility function of Player 1,

$$u(\mathbf{x}, \mathbf{y}) := \sum_{j=1}^n a^j \cdot l(x^j - y^j),$$

captures the total excess of the first army over the second army. The result on a battlefield j is $a^j \cdot l(x^j - y^j)$, where $a^j > 0$ is a weight of battlefield j and $l(x^j - y^j)$ measures the performance of the first army on a battlefield j . The standard choice is the signum function $l(z) = \text{sgn}(z)$; see (Gross and Wagner 1950) or (Roberson 2006). In this paper we assume that each player must allocate a sufficiently higher proportion of forces than the opponent to win the battle on a single battlefield. Namely, we consider

$$l(z) = \begin{cases} -1 & \text{if } z \leq -c, \\ \frac{1}{c}z & \text{if } z \in [-c, c], \\ 1 & \text{if } z \geq c, \end{cases} \quad \text{for some } c > 0. \quad (8)$$

When $c \rightarrow 0$, we recover the classical infinite colonel Blotto game since (8) approaches $\text{sgn}(z)$ in the limit.

We will show how to compute best response strategies in case of (8). Assume that Player 2 employs strategies $(\mathbf{y}_1, \dots, \mathbf{y}_k)$ with probabilities (q_1, \dots, q_k) , where $\mathbf{y}_i := (y_i^1, \dots, y_i^n) \in Y$. Then any best response strategy of Player 1 is a solution to

$$\max_{\mathbf{x} \in X} \sum_{i=1}^k q_i \sum_{j=1}^n a^j \cdot l(x^j - y_i^j). \quad (9)$$

Since l is a piecewise affine function, this nonlinear optimization problem can be reformulated as a mixed-integer linear problem. In Appendix C we derive its equivalent form

$$\begin{aligned} \max_{\mathbf{x}, \mathbf{s}, \mathbf{t}, \mathbf{z}, \mathbf{w}} \quad & \sum_{i=1}^k q_i \sum_{j=1}^n a^j (s_{ij} - t_{ij} - 1) \\ \text{s.t.} \quad & \mathbf{x} \in X, \\ & s_{ij} \geq 0, \quad s_{ij} \geq \frac{1}{c}(x^j - y_i^j + c), \\ & s_{ij} \leq \frac{1}{c}(x^j - y_i^j + c) + M_l^s(1 - z_{ij}), \\ & s_{ij} \leq M_u^s z_{ij}, \\ & t_{ij} \geq 0, \quad t_{ij} \geq \frac{1}{c}(x^j - y_i^j - c), \\ & t_{ij} \leq \frac{1}{c}(x^j - y_i^j - c) + M_l^t(1 - w_{ij}), \\ & t_{ij} \leq M_u^t w_{ij}, \\ & s_{ij} \in \mathbb{R}, \quad t_{ij} \in \mathbb{R}, \quad z_{ij} \in \{0, 1\}, \quad w_{ij} \in \{0, 1\}, \end{aligned}$$

where $M_l^s = M_u^s = \frac{1}{c} - 1$ and $M_l^t = M_u^t = \frac{1}{c} + 1$. The best response of Player 2 is obtained by solving an analogous MILP. Note that the MILP defined above is necessarily different from the one formulated in (Ganzfried 2020).

For the numerical results we consider three battlefields ($n = 3$) with equal weights ($\mathbf{a} = (1, 1, 1)$). We observed that the choice of initial strategy sets X_1 and Y_1 is crucial. Indeed, setting

$$X_1 = Y_1 = \{(1, 0, 0), (0, 1, 0), (0, 0, 1)\}$$

provides much faster convergence than starting from a random point. The reason lies in the left-hand side of Figure 5, which shows the optimal solution produced by the double oracle algorithm for $c = \frac{1}{32}$. The optimal strategies are equidistant on the grid with distance c . This is a sensible result as the best response of Player 1 to the strategy (y_1, y_2, y_3) of Player 2 is $(y_1 + c, y_2 + c, y_3 - 2c)$. Since X_1 and Y_1 already belong to the grid, all the iterates stay in it. However, they may not converge within this set when initial strategies are chosen at random.

The previous observation inspired us to start with both X_1 and Y_1 as the whole grid. It turned out that the double oracle converged in one iteration (the initial point was already an equilibrium) to the strategies depicted in Figure 6. The left-hand side shows the results for $c = \frac{1}{16}$, while the right-hand side corresponds to $c = \frac{1}{32}$. These results are close to the hexagonal solutions obtained in (Gross and Wagner 1950) and (Roberson 2006).




Figure 5: The optimal strategy for $c = \frac{1}{32}$ when started from three corner points (left). The convergence of the double oracle algorithm for $n = 3$ and $n = 10$ (scaled by $\frac{1}{50}$ for demonstration purposes) battlefields (right).

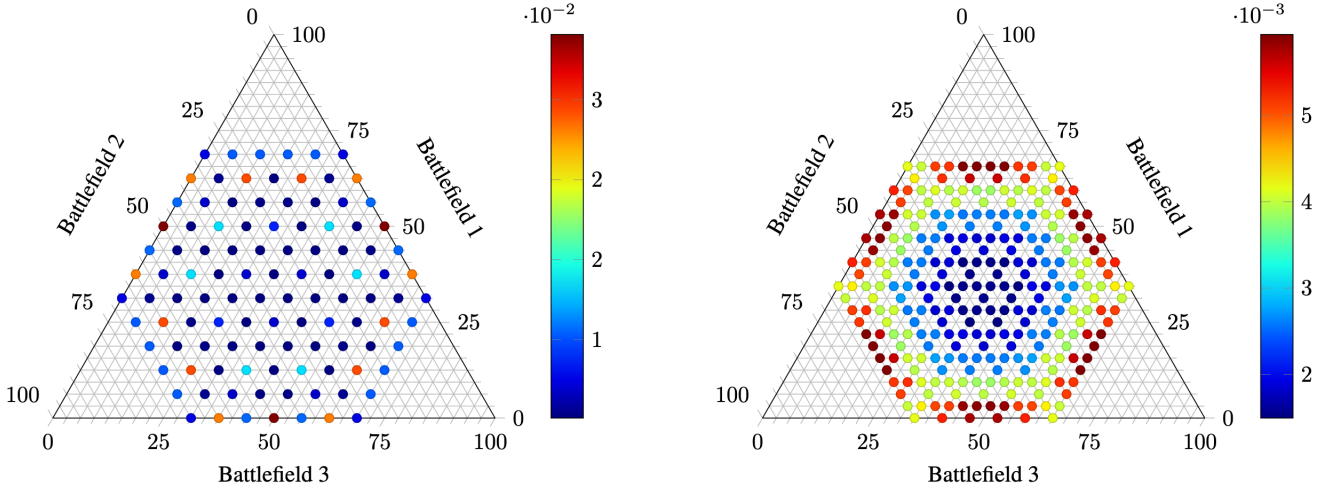


Figure 6: The optimal strategies for $c = \frac{1}{16}$ (left) and $c = \frac{1}{32}$ (right) produced by the double oracle algorithm when started from the grid. Both solutions are symmetric.

The right-hand side of Figure 5 shows the convergence of the double oracle algorithm for $n = 3$ with $\mathbf{a} = (1, 1, 1)$ and for $n = 10$ with $\mathbf{a} = (3, 4, \dots, 12)$. In both cases we put $c = \frac{1}{16}$. It appears that the convergence is influenced by c more than by the number of battlefields n .

5 Conclusions

We extended the double oracle algorithm from finite to continuous games. We proved that the algorithm recovers a finitely-supported ϵ -equilibrium in finitely many iterations and converges to an equilibrium in a possibly infinite number of iterations. We showed that the double oracle algorithm performs better than fictitious play on selected examples. It is evident that the convergence of this algorithm depends

on the size of constructed subgames and the best response calculation in each iteration. One of the open problems for future research is to analyze the speed of convergence of the double oracle algorithm.

Acknowledgments

This work was supported by the project RCI (CZ.02.1.01/0.0/0.0/16 019/0000765) “Research Center for Informatics”. T. Kasl and T. Kroupa also gratefully acknowledge the grant “Game Theory for Adversarial Machine Learning” funded by The Army Research Office.

A Weak Convergence of Measures

We will summarize a necessary background in weak topology on the space of probability measures (Billingsley 1968). A sequence of mixed strategies (p_i) in Δ_X *weakly converges* to $p \in \Delta_X$ if

$$\lim_{i \rightarrow \infty} \int_X f(x) dp_i = \int_X f(x) dp$$

for every continuous function $f: X \rightarrow \mathbb{R}$, and we denote this by $p_i \Rightarrow p$. Endowed with the topology corresponding to weak convergence, the convex set of mixed strategies Δ_X is a compact space. Analogously, Δ_Y becomes a compact set and so is the set $\Delta = \Delta_X \times \Delta_Y$. Then the definition (1) warrants that U is a continuous function on Δ . Note that compactness of Δ and continuity of U imply the existence of all maximizers/minimizers throughout the paper.

Proposition A.1. *The space Δ is weakly sequentially compact, that is, every sequence in Δ contains a weakly convergent subsequence.*

Since U is continuous, the definition of weak convergence immediately implies the following two statements:

- If $p_i \Rightarrow p$ in Δ_X and $q_i \Rightarrow q$ in Δ_Y , then

$$U(p_i, q_i) \rightarrow U(p, q). \quad (10)$$

- If $p_i \Rightarrow p$ in Δ_X and $y_i \rightarrow y$ in Y , then

$$U(p_i, y_i) \rightarrow U(p, y). \quad (11)$$

Finally, it can be shown that the optimal value of utility function in response to the opponent's mixed strategy is attained for some pure strategy.

Proposition A.2. *For any $p \in \Delta_X$ we have*

$$\min_{y \in Y} U(p, y) = \min_{q \in \Delta_Y} U(p, q).$$

B Proofs and Additional Results

Proof of Proposition 2.2. The existence of an ϵ -equilibrium follows from Theorem 3.1. To prove the second part, assume that (p^*, q^*) is an ϵ -equilibrium. Then (3) implies

$$\max_{p \in \Delta_X} U(p, q^*) - \epsilon \leq U(p^*, q^*) \leq \min_{q \in \Delta_Y} U(p^*, q) + \epsilon. \quad (12)$$

Let (\hat{p}, \hat{q}) be an equilibrium of \mathcal{G} . Then

$$U(\hat{p}, \hat{q}) \leq U(\hat{p}, q^*) \leq \max_{p \in \Delta_X} U(p, q^*) \leq U(p^*, q^*) + \epsilon,$$

where the first inequality follows from (2) and the third from (12). In a similar way, we can show

$$U(\hat{p}, \hat{q}) \geq U(p^*, \hat{q}) \geq \min_{q \in \Delta_Y} U(p^*, q) \geq U(p^*, q^*) - \epsilon.$$

Combining these two relations with $U(\hat{p}, \hat{q}) = v(\mathcal{G})$ imply the second statement of Proposition 2.2. \square

Lemma B.1. *Assume $X_{i+1} = X_i$ and $Y_{i+1} = Y_i$ in some step i of Algorithm 3.1. Then $U(p_i^*, y_{i+1}) = U(x_{i+1}, q_i^*)$.*

Proof. The condition $X_{i+1} = X_i$ implies $x_{i+1} \in X_i$. Then $U(p_i^*, q_i^*) = \max_{x \in X_i} U(x, q_i^*) = \max_{x \in X} U(x, q_i^*) = U(x_{i+1}, q_i^*)$,

where the first equality follows from Proposition 2.1 applied to the subgame (X_i, Y_i, u) , the second from $x_{i+1} \in X_i$, and the third from the definition of iterate x_{i+1} .

Similarly, we can show $U(p_i^*, q_i^*) = U(p_i^*, y_{i+1})$, which means $U(p_i^*, y_{i+1}) = U(x_{i+1}, q_i^*)$. \square

Lemma B.2. *The inequality*

$$v_i \leq v(\mathcal{G}) \leq \bar{v}_i$$

holds in every step i of Algorithm 3.1.

Proof. Let (p^*, q^*) be an equilibrium of \mathcal{G} . Then

$$\begin{aligned} v_i &= U(p_i^*, y_{i+1}) = \min_{y \in Y} U(p_i^*, y) = \min_{q \in \Delta_Y} U(p_i^*, q) \\ &\leq U(p_i^*, q^*) \leq U(p^*, q^*) = v(\mathcal{G}). \end{aligned}$$

The second inequality can be obtained analogously. \square

Example B.1. Define $X := [0, 1]$, $Y := [0, 1]$, and consider an arbitrary continuous function $u: X \times Y \rightarrow \mathbb{R}$ for which the double oracle algorithm produces an infinite number of iterates $(x_1, y_1), (x_2, y_2), \dots$ for $\epsilon = 0$. Further, put $\tilde{X} := [0, 1] \cup [2, 3]$ and let $\tilde{u}: \tilde{X} \times Y \rightarrow \mathbb{R}$ be given by

$$\tilde{u}(x, y) = \begin{cases} u(x, y) & \text{if } x \in [0, 1], \\ u(x-2, y) & \text{if } x \in [2, 3]. \end{cases}$$

Since u is continuous, $(\tilde{X}, Y, \tilde{u})$ is a continuous game. Since $\tilde{u}(x, y) = \tilde{u}(x+2, y)$, the extrema of marginal functions are not unique. Considering $\tilde{y}_i = y_i$, the double oracle algorithm may produce the sequence of iterations

$$\tilde{x}_i = \begin{cases} x_i & \text{if } i \text{ is odd,} \\ x_i + 2 & \text{if } i \text{ is even.} \end{cases}$$

This sequence is obviously not convergent. However, there exists a convergent subsequence and its limit is an equilibrium by Theorem 3.1.

C Best Response for Colonel Blotto Game

Function l from (8) can be written as

$$l(z) = \max \left\{ \frac{1}{c}(z+c), 0 \right\} - \max \left\{ \frac{1}{c}(z-c), 0 \right\} - 1.$$

With each i, j in (9) we associate auxiliary variables s_{ij} and t_{ij} and the constraints ensuring $l(x^j - y_i^j) = s_{ij} - t_{ij} - 1$. The constraints on s_{ij} and t_{ij} follow from Lemma C.1.

Lemma C.1. *Let $a > 0$, $b \in \mathbb{R}$, $M_l > 0$, $M_u > 0$ and $f(x) := \max\{a(x-b), 0\}$. For every x such that $a(x-b) \in [-M_l, M_u]$ there are a unique $s \in \mathbb{R}$ and a possibly non-unique $z \in \{0, 1\}$ solving the system*

$$\begin{aligned} s &\geq 0, & s &\leq a(x-b) + M_l(1-z), \\ s &\geq a(x-b), & s &\leq M_u z. \end{aligned}$$

Moreover, it holds $f(x) = s$.

Proof. The proof is based on the well-known big-M method for the deactivation of constraints. The claim follows from the following implications,

$$\begin{aligned} a(x - b) < 0 &\implies z = 0 \implies s = 0, \\ a(x - b) > 0 &\implies z = 1 \implies s = a(x - b). \end{aligned}$$

If $a(x - b) = 0$, then $s = 0$ is unique, whereas z may have either value. \square

Since $x^j, y_i^j \in [0, 1]$, we have

$$\begin{aligned} \frac{1}{c}(x^j - y_i^j + c) &\in [-\frac{1}{c} + 1, \frac{1}{c} + 1], \\ \frac{1}{c}(x^j - y_i^j - c) &\in [-\frac{1}{c} - 1, \frac{1}{c} - 1], \end{aligned}$$

which gives the bounds in Lemma C.1.

References

- Archibald, C.; and Shoham, Y. 2009. Modeling billiards games. In *Proceedings of the 8th International Conference on Autonomous Agents and Multiagent Systems*, volume 1, 193–199.
- Behnezhad, S.; Dehghani, S.; Derakhshan, M.; HajiAghayi, M.; and Seddighin, S. 2017. Faster and simpler algorithm for optimal strategies of Blotto game. In *Proceedings of the 31st AAAI Conference on Artificial Intelligence*, 369–375.
- Billingsley, P. 1968. *Convergence of probability measures*. New York: John Wiley & Sons Inc.
- Brown, G. W. 1951. Iterative solution of games by fictitious play. In *Activity analysis of production and allocation*, 374–376. John Wiley & Sons.
- Brown, N.; and Sandholm, T. 2019. Superhuman AI for multiplayer poker. *Science* 365(6456): 885–890.
- Fan, K. 1952. Fixed-point and minimax theorems in locally convex topological linear spaces. *Proceedings of the National Academy of Sciences of the United States of America* 38(2): 121.
- Ganzfried, S. 2020. Algorithm for Computing Approximate Nash equilibrium in Continuous Games with Application to Continuous Blotto. *arXiv preprint arXiv:2006.07443*.
- Gilpin, A.; Peña, J.; and Sandholm, T. 2012. First-order algorithm with $\mathcal{O}(\ln(1/\epsilon))$ convergence for ϵ -equilibrium in two-person zero-sum games. *Mathematical Programming* 133(1-2): 279–298.
- Glicksberg, I. L. 1952. A further generalization of the Kakutani fixed point theorem, with application to Nash equilibrium points. *Proceedings of the American Mathematical Society* 3: 170–174.
- Gross, O.; and Wagner, R. 1950. A continuous Colonel Blotto game. Technical report, RM-408.
- Hofbauer, J.; and Sorin, S. 2006. Best response dynamics for continuous zero-sum games. *Discrete and Continuous Dynamical Systems–Series B* 6(1): 215.
- Kamra, N.; Gupta, U.; Fang, F.; Liu, Y.; and Tambe, M. 2018. Policy learning for continuous space security games using neural networks. In *Thirty-Second AAAI Conference on Artificial Intelligence*, 1103–1112.
- Kamra, N.; Gupta, U.; Wang, K.; Fang, F.; Liu, Y.; and Tambe, M. 2019. DeepFP for Finding Nash Equilibrium in Continuous Action Spaces. In *International Conference on Decision and Game Theory for Security*, 238–258. Springer.
- Karlin, S. 1959. *Mathematical Methods and Theory in Games, Programming and Economics*. Vol. 2: *The Theory of Infinite Games*. Addison-Wesley Publishing Company.
- Lanctot, M.; Zambaldi, V. F.; Gruslys, A.; Lazaridou, A.; Tuyls, K.; Pérolat, J.; Silver, D.; and Graepel, T. 2017. A Unified Game-Theoretic Approach to Multiagent Reinforcement Learning. In *Advances in Neural Information Processing Systems 30: Annual Conference on Neural Information Processing Systems*, 4190–4203.
- Laraki, R.; and Lasserre, J. B. 2012. Semidefinite programming for min–max problems and games. *Mathematical programming* 131(1-2): 305–332.
- McMahan, H. B.; Gordon, G. J.; and Blum, A. 2003. Planning in the presence of cost functions controlled by an adversary. In *Proceedings of the 20th International Conference on Machine Learning (ICML-03)*, 536–543.
- Moravčík, M.; Schmid, M.; Burch, N.; Lisý, V.; Morrill, D.; Bard, N.; Davis, T.; Waugh, K.; Johanson, M.; and Bowling, M. 2017. Deepstack: Expert-level artificial intelligence in heads-up no-limit poker. *Science* 356(6337): 508–513.
- Muller, P.; Omidshafiei, S.; Rowland, M.; Tuyls, K.; Perolat, J.; Liu, S.; Hennes, D.; Marris, L.; Lanctot, M.; Hughes, E.; et al. 2019. A generalized training approach for multiagent learning. *arXiv preprint arXiv:1909.12823*.
- Parrilo, P. 2006. Polynomial games and sum of squares optimization. In *Decision and Control, 2006 45th IEEE Conference on*, 2855–2860.
- Perkins, S.; and Leslie, D. S. 2014. Stochastic fictitious play with continuous action sets. *Journal of Economic Theory* 152: 179–213.
- Rehbeck, J. 2018. Note on unique Nash equilibrium in continuous games. *Games and Economic Behavior* 110: 216–225.
- Roberson, B. 2006. The Colonel Blotto game. *Economic Theory* 29: 1–24.
- Robinson, J. 1951. An iterative method of solving a game. *Annals of Mathematics* (54): 296–301.
- Stein, N. D.; Ozdaglar, A.; and Parrilo, P. A. 2008. Separable and low-rank continuous games. *International Journal of Game Theory* 37(4): 475–504.
- Townsend, A. 2014. Constrained optimization in Chebfun. URL <http://www.chebfun.org/examples/opt/ConstrainedOptimization.html>.
- Yasodharan, S.; and Loiseau, P. 2019. Nonzero-sum Adversarial Hypothesis Testing Games. In *Advances in Neural Information Processing Systems*, 7310–7320.